HTTP Key Value Server

—

Dev Sonar 25M0824

System Architecture

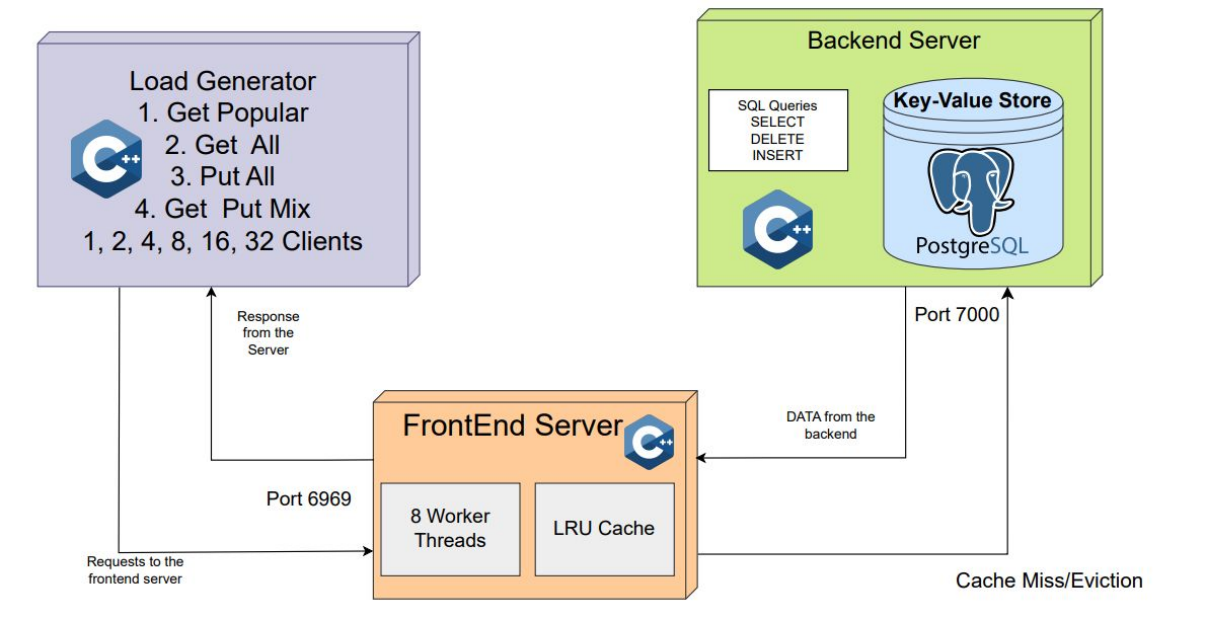

a

Sequence Diagram (Request Flow)

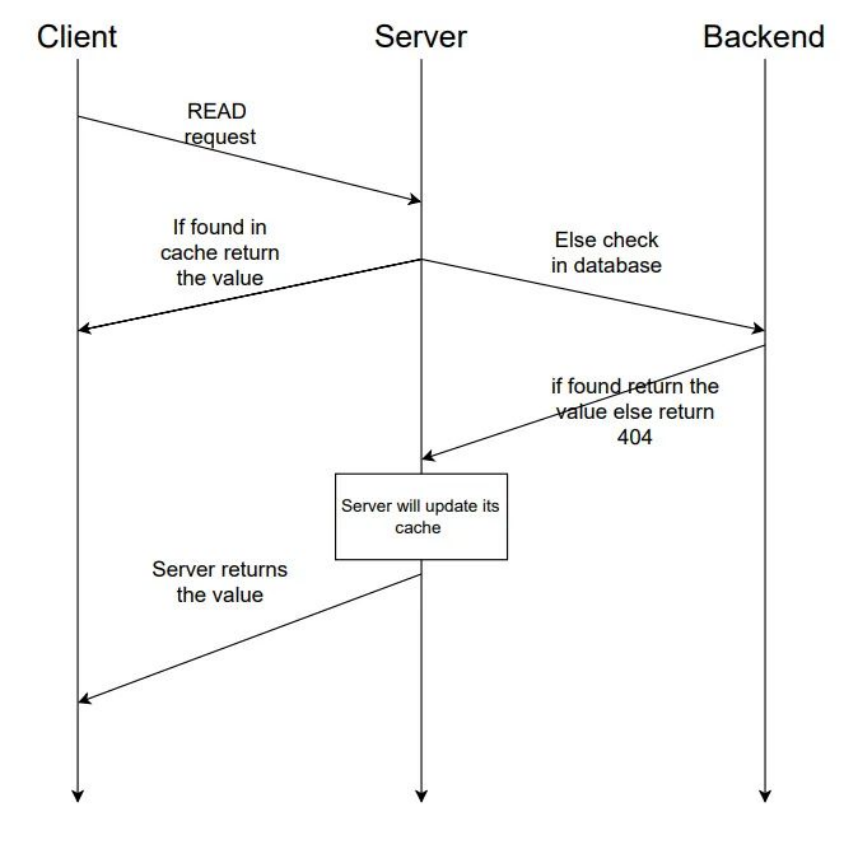

Load Generator Architecture

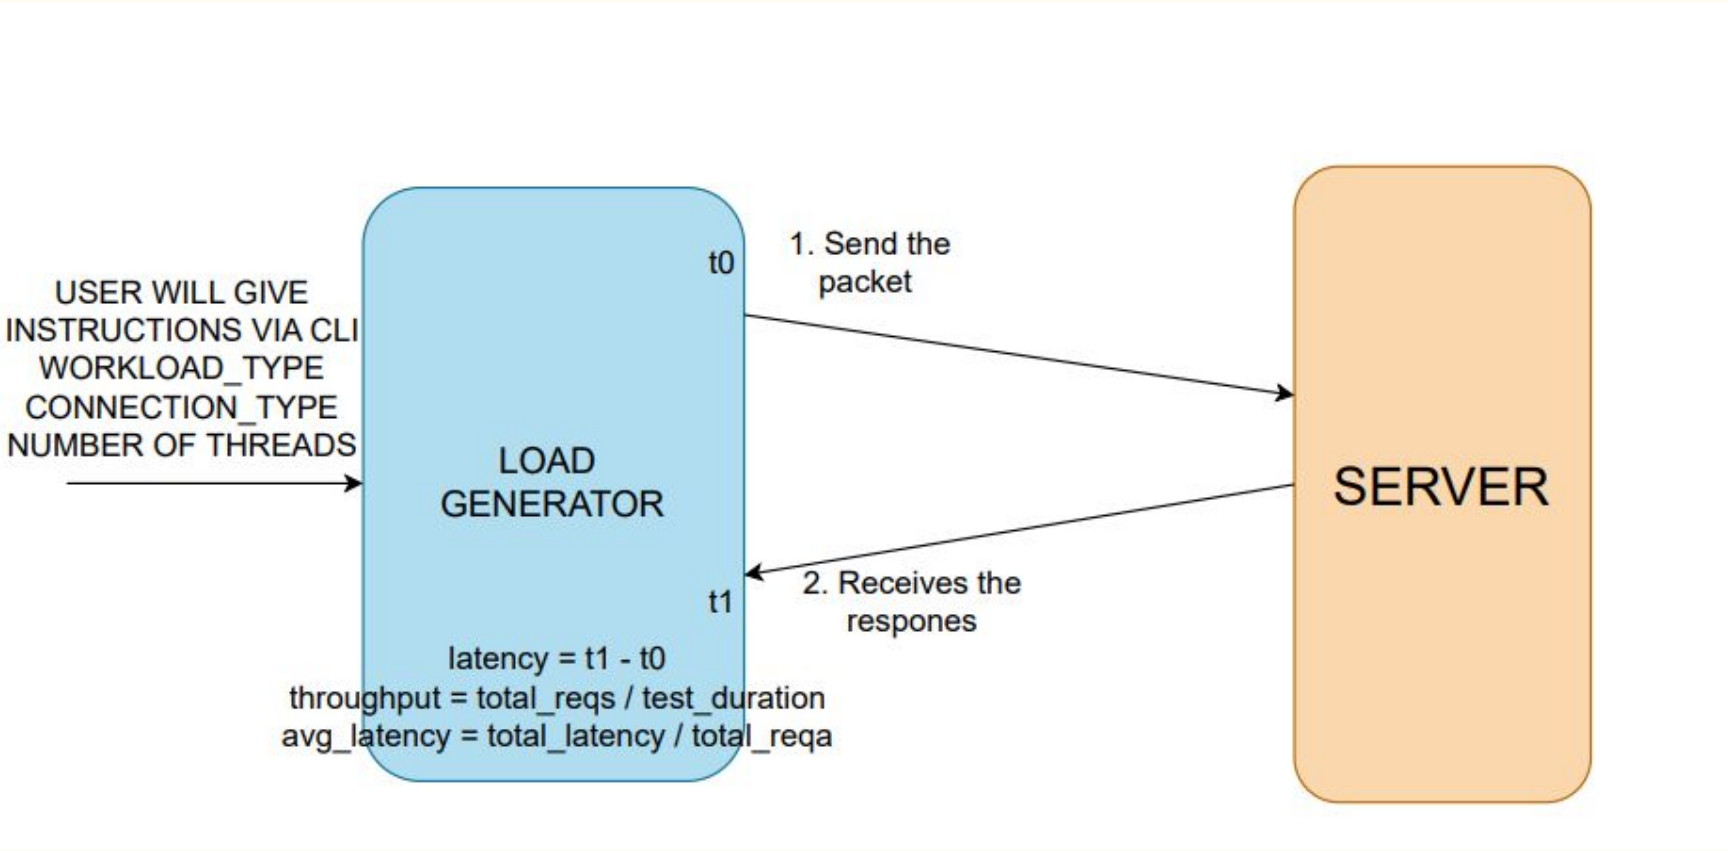

Load Test Setup

1. Everything was tested locally
2. Cloud methods were tried but they were not effective enough
3. The backend_server was pinned to a single core number : 3
4. The frontend_server was pinned to cores 4-7
5. The load_generator was pinned to cores 1-2
6. The bottlenecks for CPU and IO were seen using the htop and iostat command respectively
7. Processor : i5-10300H

CPU Bottleneck

It was achieved while using the GET_POPULAR workload

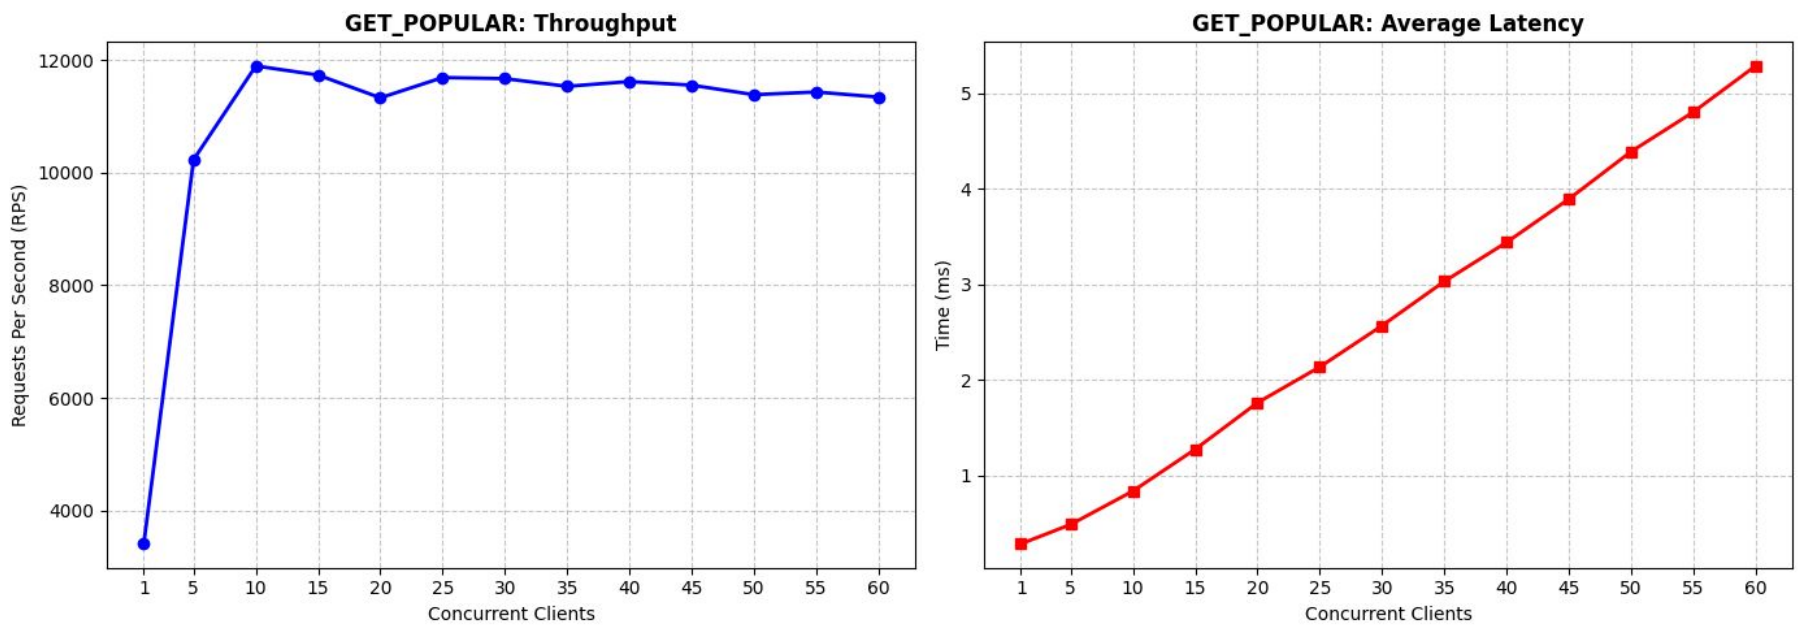

IO Bottleneck

It was achieved while using the GET_ALL workload

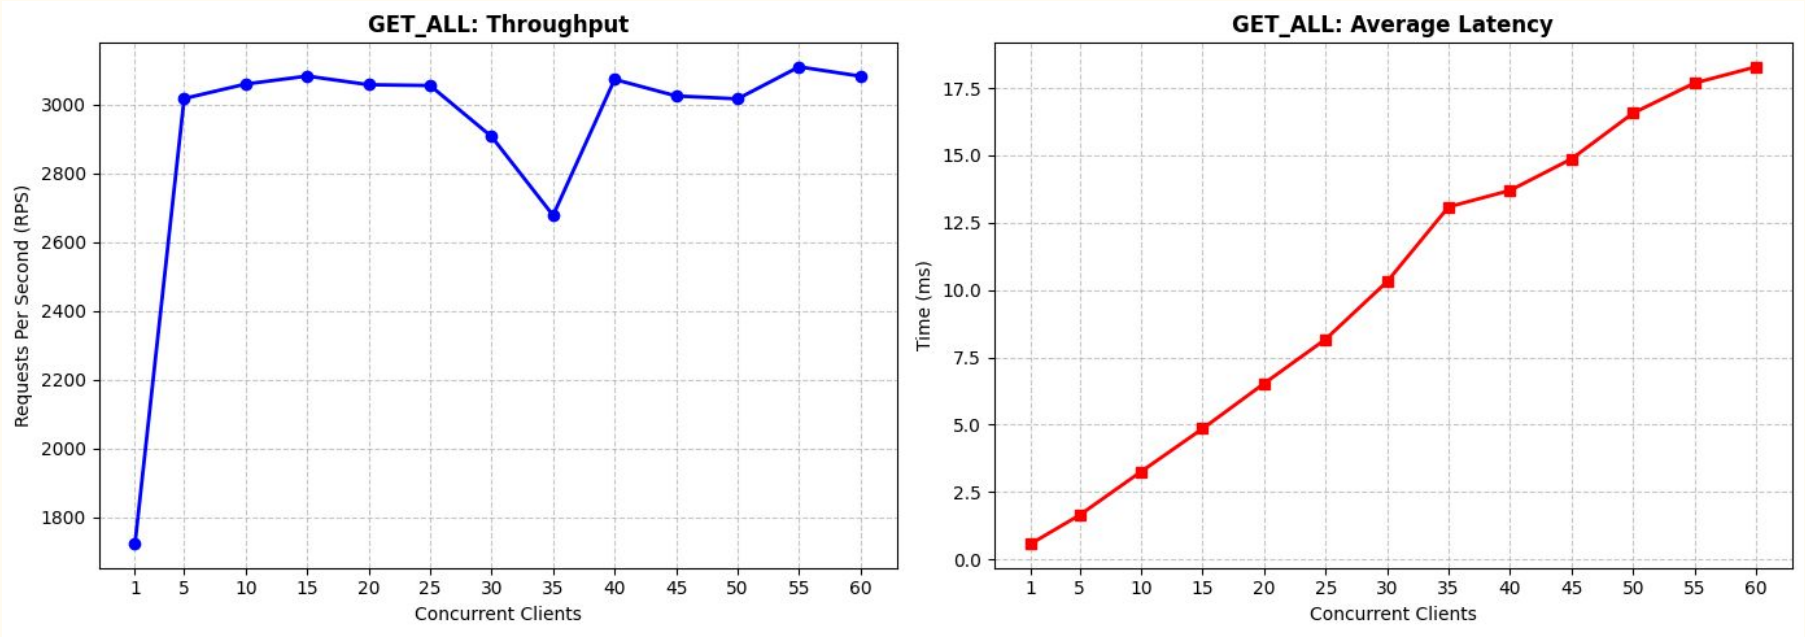

Other workloads :

PUT_ALL

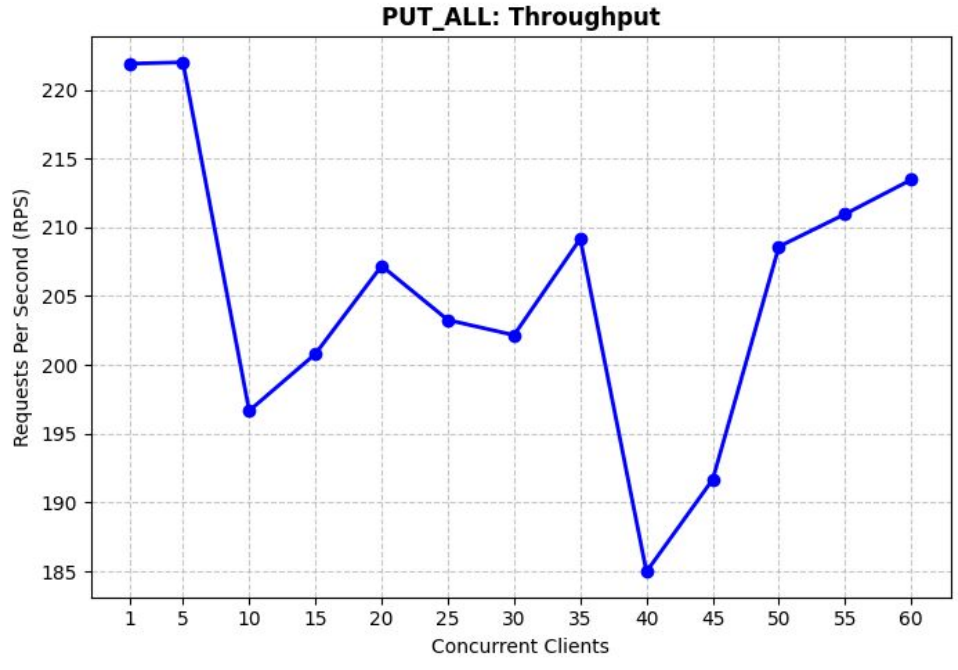

GET_PUT_MIX

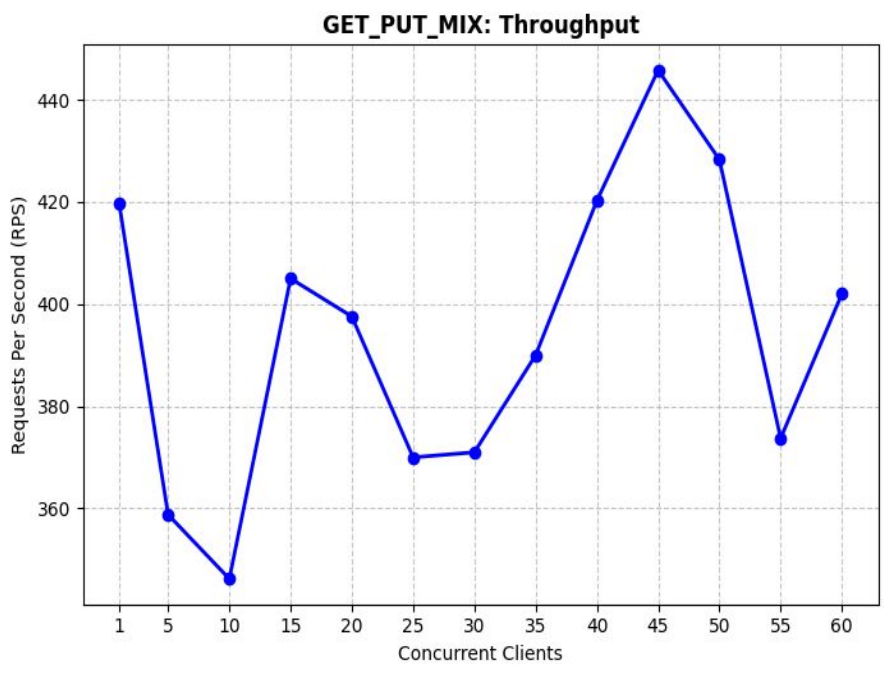
GET_PUT_MIX: Throughput
Requests Per Second (RPS)
360
380
400
420
440
1
5
10
15
20
25
30
35
40
45
50
55
60
Concurrent Clients

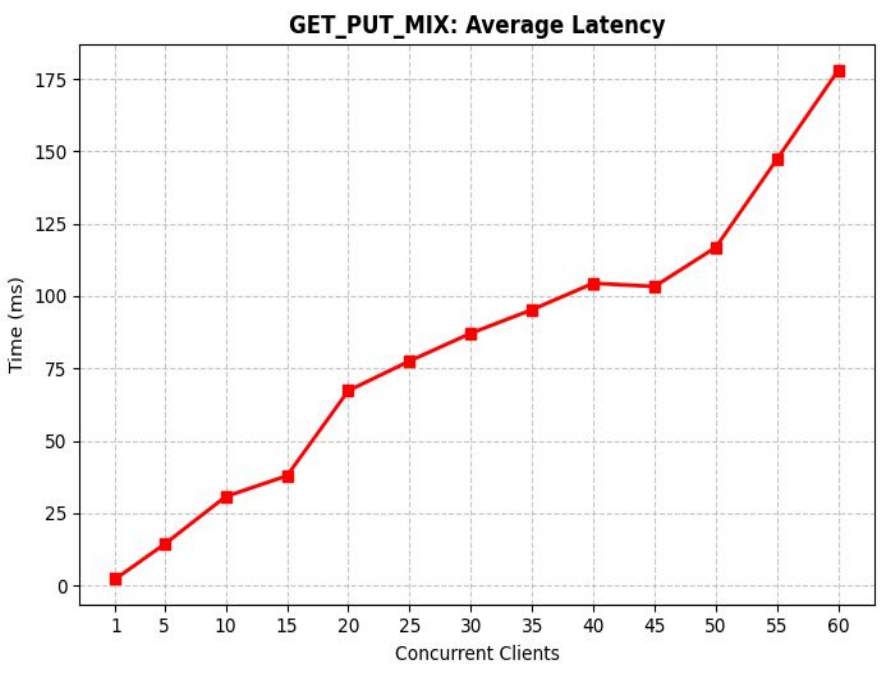
GET_PUT_MIX: Average Latency
Time (ms)
0
25
50
75
100
125
150
175
1
5
10
15
20
25
30
35
40
45
50
55
60
Concurrent Clients

Some snapshots of performance testing :

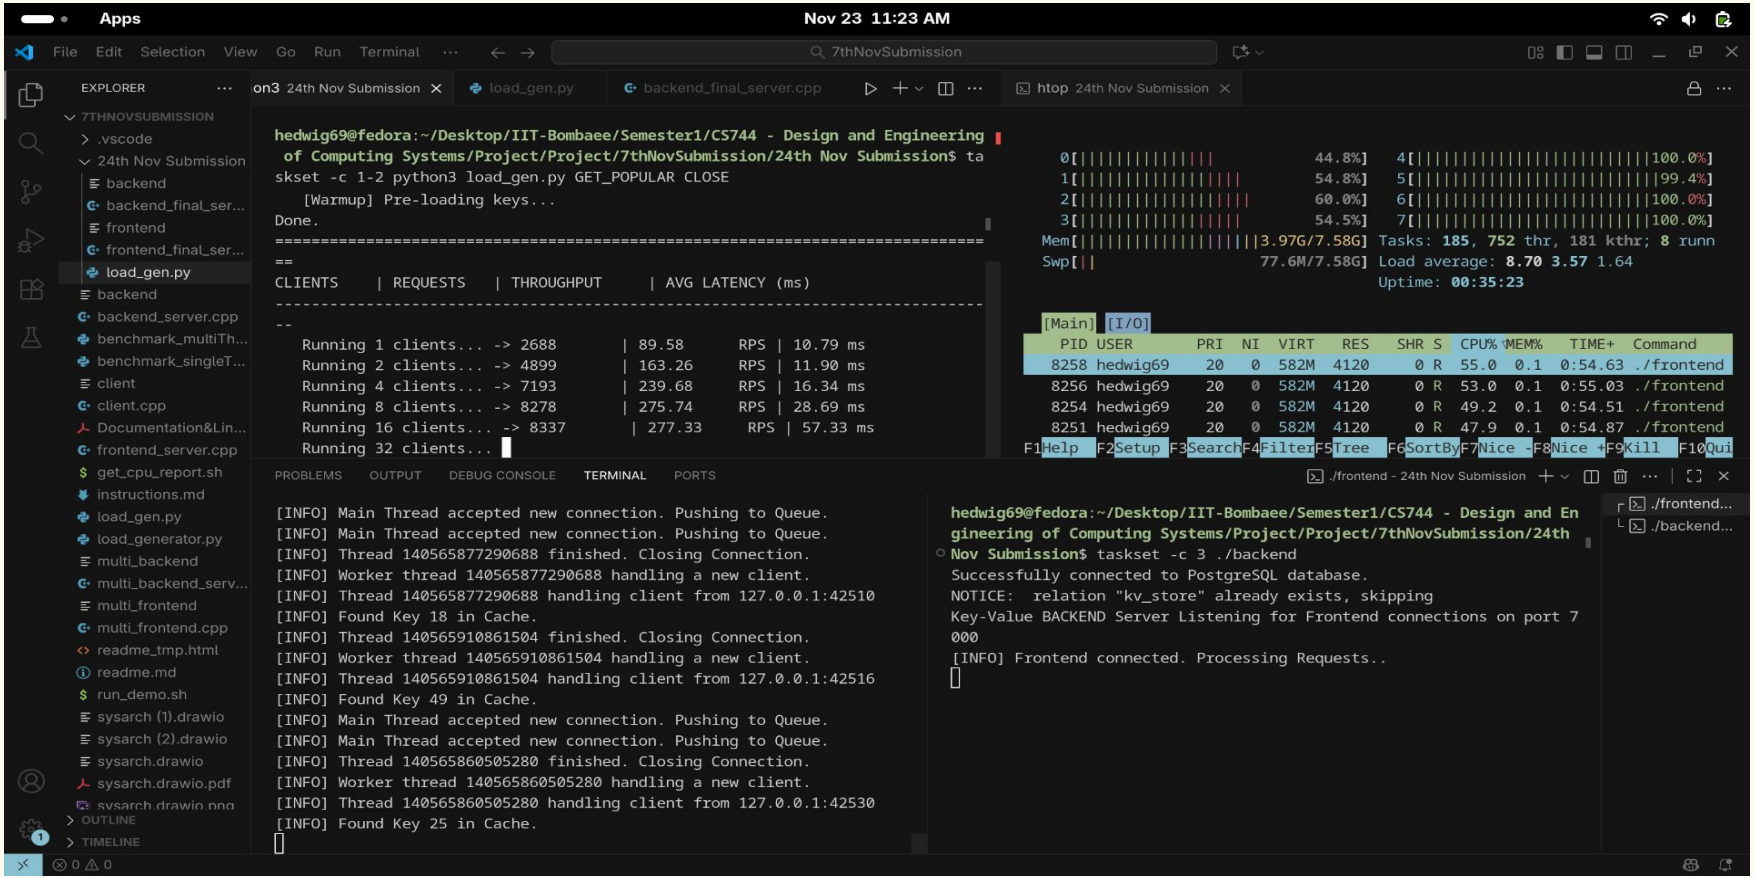

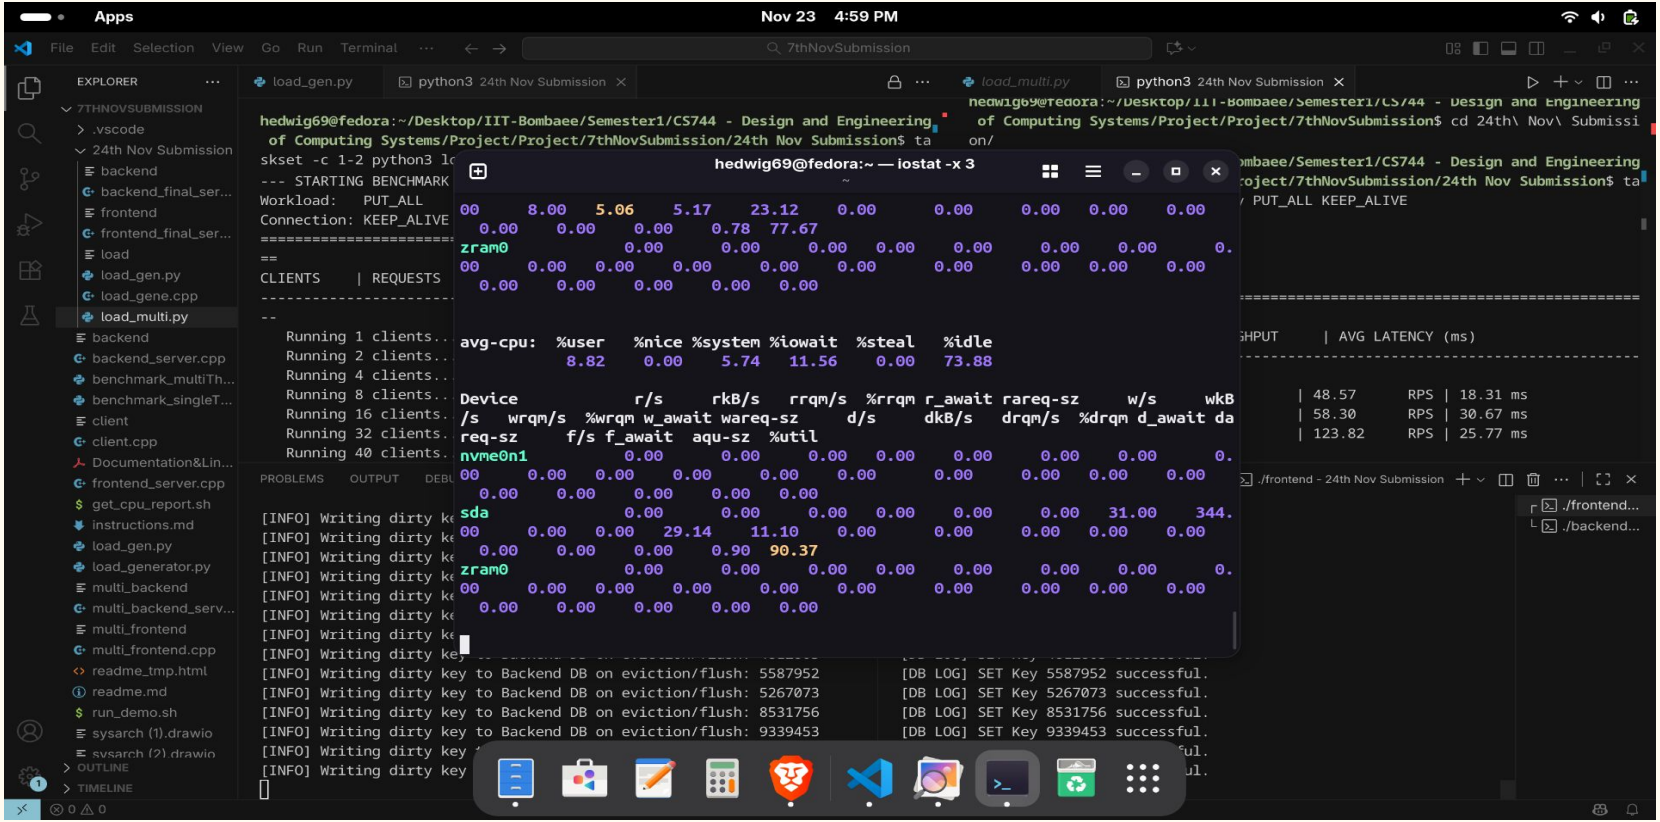
Apps
Nov 23 4:59 PM
hedwig69@fedora:~ — iostat -x 3
avg-cpu: %user %nice %system %iowait %steal %idle
8.82 0.00 5.74 11.56 0.00 73.88
Device r/s rkB/s rrqm/s %rrqm r_await rareq-sz w/s wkB/s wrqm/s %wrqm w_await wareq-sz d/s dkB/s drqm/s %drqm d_await dareq-sz f/s f_await aqu-sz %util
nvme0n1 0.00 0.00 0.00 0.00 0.00 0.00 0.00 0.00 0.00 0.00 0.00 0.00 0.00 0.00 0.00 0.00 0.00 0.00
sda 0.00 0.00 0.00 0.00 0.00 0.00 0.00 0.00 31.00 344.00 29.14 11.10 0.00 0.00 0.00 0.00 0.90 90.37
zram0 0.00 0.00 0.00 0.00 0.00 0.00 0.00 0.00 0.00 0.00 0.00 0.00 0.00 0.00 0.00 0.00 0.00 0.00
[INFO] Writing dirty key to Backend DB on eviction/flush: 5587952
[DB LOG] SET Key 5587952 successful.
[INFO] Writing dirty key to Backend DB on eviction/flush: 5267073
[DB LOG] SET Key 5267073 successful.
[INFO] Writing dirty key to Backend DB on eviction/flush: 8531756
[DB LOG] SET Key 8531756 successful.
[INFO] Writing dirty key to Backend DB on eviction/flush: 9339453
[DB LOG] SET Key 9339453 successful.